

Task_4_Churn_Prediction_Model

November 24, 2025

1 1. Load and Clean the Dataset

- 1.1 Import Libraries
- 1.2 Data From GitHub
- 1.3 Check the missing values
- 1.4 Convert TotalCharges to Numeric
- 1.5 Fill the missing value with median
- 1.6 Also Check the " " value
- 1.7 Convert types

```
[73]: # 1.1 Import Libraries
import pandas as pd
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report, confusion_matrix
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

import matplotlib.pyplot as plt
import seaborn as sns
```

```
[74]: # 1.2 Data From GitHub

# GitHub raw URL
url = 'https://raw.githubusercontent.com/abuthahir17/Dataset/main/Telco_Customer_Churn_Dataset.csv'

# Read CSV file
data = pd.read_csv(url)

print("Dataset Loaded Successfully!")
```

Dataset Loaded Successfully!

```
[75]: # 1.3 Check the missing values  
data.replace(" ", None, inplace=True)  
print("Missing Value: \n", data.isnull().sum())
```

```
Missing Value:  
customerID      0  
gender          0  
SeniorCitizen   0  
Partner         0  
Dependents     0  
tenure          0  
PhoneService    0  
MultipleLines   0  
InternetService 0  
OnlineSecurity  0  
OnlineBackup    0  
DeviceProtection 0  
TechSupport     0  
StreamingTV    0  
StreamingMovies 0  
Contract        0  
PaperlessBilling 0  
PaymentMethod   0  
MonthlyCharges  0  
TotalCharges    11  
Churn           0  
dtype: int64
```

```
[76]: # 1.4 Convert TotalCharges to Numeric  
data["TotalCharges"] = pd.to_numeric(data["TotalCharges"], errors="coerce")
```

```
[77]: # 1.5 Fill the missing value with median  
data["TotalCharges"] = data["TotalCharges"].fillna(data["TotalCharges"].  
median())
```

```
[78]: # 1.6 Also Check the " " value  
data.replace(" ", None, inplace=True)  
print("Missing Value after Cleaning: \n", data.isnull().sum())
```

```
Missing Value after Cleaning:  
customerID      0  
gender          0  
SeniorCitizen   0  
Partner         0  
Dependents     0  
tenure          0  
PhoneService    0  
MultipleLines   0
```

```
InternetService      0
OnlineSecurity       0
OnlineBackup         0
DeviceProtection     0
TechSupport          0
StreamingTV          0
StreamingMovies       0
Contract             0
PaperlessBilling      0
PaymentMethod         0
MonthlyCharges        0
TotalCharges          0
Churn                0
dtype: int64
```

```
[79]: # 1.7 Convert types
data["Churn"] = data["Churn"].map({"Yes":1, "No":0})
```

2 2. Encoding Categorical Variables

- 2.1 Identify Numerical and Categorical Columns
- 2.2 Before Encoding
- 2.3 During Encoding
- 2.4 After Encoding

```
[80]: # 2.1 Identify numerical columns
numerical_cols = data.select_dtypes(include=['int64', 'float64']).columns
print("Total number of Numerical Columns:", len(numerical_cols))
print("All Numerical Columns:" ,list(numerical_cols), "\n")

# 2.1 Identify categorical columns
categorical_cols = data.select_dtypes(include=['object']).columns
print("Total number of Categorical Columns:", len(categorical_cols))
print("All Categorical Columns:" , list(categorical_cols))
```

```
Total number of Numerical Columns: 5
All Numerical Columns: ['SeniorCitizen', 'tenure', 'MonthlyCharges',
'TotalCharges', 'Churn']
```

```
Total number of Categorical Columns: 16
All Categorical Columns: ['customerID', 'gender', 'Partner', 'Dependents',
'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod']
```

```
[81]: # 2.2 Before Encoding
print("Before encoding:\n", data.head())
```

```
print("Shape before encoding:", data.shape)
```

Before encoding:

```
customerID gender SeniorCitizen Partner Dependents tenure PhoneService \
0 7590-VHVEG Female 0 Yes No 1 No
1 5575-GNVDE Male 0 No No 34 Yes
2 3668-QPYBK Male 0 No No 2 Yes
3 7795-CFOCW Male 0 No No 45 No
4 9237-HQITU Female 0 No No 2 Yes

MultipleLines InternetService OnlineSecurity ... DeviceProtection \
0 No phone service DSL No ...
1 No DSL Yes ...
2 No DSL Yes ...
3 No phone service DSL Yes ...
4 No Fiber optic No ...

TechSupport StreamingTV StreamingMovies Contract PaperlessBilling \
0 No No No Month-to-month Yes
1 No No No One year No
2 No No No Month-to-month Yes
3 Yes No No One year No
4 No No No Month-to-month Yes

PaymentMethod MonthlyCharges TotalCharges Churn
0 Electronic check 29.85 29.85 0
1 Mailed check 56.95 1889.50 0
2 Mailed check 53.85 108.15 1
3 Bank transfer (automatic) 42.30 1840.75 0
4 Electronic check 70.70 151.65 1
```

[5 rows x 21 columns]

Shape before encoding: (7043, 21)

[82]: # 2.3 During Encoding

```
from sklearn.preprocessing import LabelEncoder

# Binary columns (2 unique values)
binary_cols = [col for col in categorical_cols if data[col].nunique() == 2]
print("Binary Columns (LabelEncode):", binary_cols, "\n")

# Multi-category columns (>2 unique values)
multi_cat_cols = [col for col in categorical_cols if data[col].nunique() > 2
                  ↵and col not in binary_cols]
print("Multi-category Columns (One-hot Encode):", multi_cat_cols, "\n")
```

```

label_encoder = LabelEncoder()

# Binary
for col in binary_cols:
    data[col] = label_encoder.fit_transform(data[col])

# Multi-Category (One-Hot Encoding)
data = pd.get_dummies(data, columns=multi_cat_cols, drop_first=True)

print("Categorical Encoding Completed.")

```

Binary Columns (LabelEncode): ['gender', 'Partner', 'Dependents',
'PhoneService', 'PaperlessBilling']

Multi-category Columns (One-hot Encode): ['customerID', 'MultipleLines',
'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract', 'PaymentMethod']

Categorical Encoding Completed.

[83]: # 2.4 After Encoding

```

print("After encoding:\n", data.head())
print("New Shape After Encoding:", data.shape)

```

After encoding:

	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	\
0	0	0	1	0	1	0	
1	1	0	0	0	34	1	
2	1	0	0	0	2	1	
3	1	0	0	0	45	0	
4	0	0	0	0	2	1	

	PaperlessBilling	MonthlyCharges	TotalCharges	Churn	...	\
0	1	29.85	29.85	0	...	
1	0	56.95	1889.50	0	...	
2	1	53.85	108.15	1	...	
3	0	42.30	1840.75	0	...	
4	1	70.70	151.65	1	...	

	TechSupport_Yes	StreamingTV_No internet service	StreamingTV_Yes	\
0	False	False	False	
1	False	False	False	
2	False	False	False	
3	True	False	False	
4	False	False	False	

	StreamingMovies_No internet service	StreamingMovies_Yes	\
--	-------------------------------------	---------------------	---

```

0                         False          False
1                         False          False
2                         False          False
3                         False          False
4                         False          False

  Contract_One year  Contract_Two year \
0                 False          False
1                 True           False
2                False          False
3                 True           False
4                False          False

  PaymentMethod_Credit card (automatic)  PaymentMethod_Electronic check \
0                           False          True
1                           False         False
2                           False         False
3                           False         False
4                           False          True

  PaymentMethod_Mailed check
0                     False
1                     True
2                     True
3                    False
4                    False

[5 rows x 7073 columns]
New Shape After Encoding: (7043, 7073)

```

3 3. Dataset Splitting (Train/Test)

- 3.1 Fix the target variable
- 3.2 Split the Dataset

[84]: #3.1 Fix the target variable

```

from sklearn.model_selection import train_test_split

# Target variable
X = data.drop("Churn", axis=1)
y = data["Churn"]

```

[85]: # 3.2 Split the dataset into 2 set (80% train, 20% test)

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

```

```
print("Training set shape:", X_train.shape)
print("Testing set shape:", X_test.shape)
```

```
Training set shape: (5634, 7072)
Testing set shape: (1409, 7072)
```

4 4. Standardize Numerical Features

```
[86]: scaler = StandardScaler()
num_cols = ['tenure', 'MonthlyCharges', 'TotalCharges']

# scale only if columns exist in X (safety)
num_cols = [c for c in num_cols if c in X_train.columns]

X_train[num_cols] = scaler.fit_transform(X_train[num_cols])
X_test[num_cols] = scaler.transform(X_test[num_cols])
```

5 5. Fix class imbalance on training set using SMOTE

```
[87]: from imblearn.over_sampling import SMOTE

# Synthetic Minority Oversampling Technique (SMOTE) -> To fix class imbalance
# in your churn dataset.

smote = SMOTE(random_state=42)
X_train_sm, y_train_sm = smote.fit_resample(X_train, y_train)
print("After SMOTE ->", X_train_sm.shape, y_train_sm.sum(), "positive samples")
```

```
After SMOTE -> (8278, 7072) 4139 positive samples
```

6 6. Model Training

- 6.1 Logistic Regression
- 6.2 Decision Tree
- 6.3 Random Forest

```
[88]: # 6.1 Logistic Regression

log = LogisticRegression(max_iter=1000, random_state=42)
log.fit(X_train_sm, y_train_sm)
y_pred_log = log.predict(X_test)

print("Logistic Regression Results:")
print(classification_report(y_test, y_pred_log))
```

```
Logistic Regression Results:
precision    recall   f1-score   support


```

0	0.86	0.83	0.84	1035
1	0.57	0.62	0.59	374
accuracy			0.78	1409
macro avg	0.71	0.73	0.72	1409
weighted avg	0.78	0.78	0.78	1409

[89]: # 6.2 Decision Tree

```
dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train_sm, y_train_sm)
y_pred_tree = dt.predict(X_test)

print("Decision Tree Results:")
print(classification_report(y_test, y_pred_tree))
```

Decision Tree Results:

	precision	recall	f1-score	support
0	0.85	0.84	0.84	1035
1	0.57	0.60	0.59	374
accuracy			0.77	1409
macro avg	0.71	0.72	0.71	1409
weighted avg	0.78	0.77	0.78	1409

[90]: # 6.3 Random Forest

```
rf = RandomForestClassifier(n_estimators=200, random_state=42)
rf.fit(X_train_sm, y_train_sm)
y_pred_rf = rf.predict(X_test)

print("Random Forest Results:")
print(classification_report(y_test, y_pred_rf))
```

Random Forest Results:

	precision	recall	f1-score	support
0	0.86	0.86	0.86	1035
1	0.61	0.61	0.61	374
accuracy			0.79	1409
macro avg	0.73	0.74	0.74	1409
weighted avg	0.79	0.79	0.79	1409

7 7. Hyperparameter tuning for RandomForest (GridSearchCV)

```
[91]: # Hyperparameter Tuning (Decision Tree Example)
from sklearn.model_selection import GridSearchCV

param_grid = [
    'n_estimators': [100, 200],
    'max_depth': [5, 10, None],
    'min_samples_split': [2, 5]
]

grid = GridSearchCV(
    RandomForestClassifier(random_state=42),
    param_grid,
    scoring='f1',
    cv=3,
    n_jobs=-1,
    verbose=1
)

grid.fit(X_train_sm, y_train_sm)
best_rf = grid.best_estimator_
print("\nGridSearch best params:", grid.best_params_)
```

Fitting 3 folds for each of 12 candidates, totalling 36 fits

```
GridSearch best params: {'max_depth': None, 'min_samples_split': 2,
'n_estimators': 200}
```

8 8. Tuned Random Forest Evaluation

8.1 Classification Report of Tuned Random Forest

8.2 Confusion Matrix of Tuned Random Forest

8.3 ROC Curve and AUC Score


```
[92]: # 8.1 Classification Report of Tuned Random Forest
y_pred_best = best_rf.predict(X_test)
print("\nTuned Random Forest:\n", classification_report(y_test, y_pred_best))
```

Tuned Random Forest:

	precision	recall	f1-score	support
0	0.86	0.86	0.86	1035
1	0.61	0.61	0.61	374
accuracy			0.79	1409
macro avg	0.73	0.74	0.74	1409

```
weighted avg      0.79      0.79      0.79      1409
```


```
[93]: # 8.2 Confusion Matrix of Tuned Random Forest
cm = confusion_matrix(y_test, y_pred_best)
plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix - Tuned RF")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



```
[94]: # 8.3 ROC Curve and AUC Score
from sklearn.metrics import roc_curve, roc_auc_score

y_prob = best_rf.predict_proba(X_test)[:, 1]
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
auc = roc_auc_score(y_test, y_prob)
plt.figure(figsize=(6,4))
plt.plot(fpr, tpr, label=f"AUC = {auc:.3f}")
plt.plot([0,1],[0,1], 'k--')
plt.title("ROC Curve - Tuned RF")
plt.xlabel("False Positive Rate")
```


```
plt.ylabel("True Positive Rate")
plt.legend()
plt.show()
```



9 9. Feature Importance

```
[95]: # Feature Importance (Random Forest)

importances = pd.Series(best_rf.feature_importances_, index=X.columns)
top_imp = importances.sort_values(ascending=False).head(15)
plt.figure(figsize=(10,8))
top_imp.plot(kind='barh')
plt.gca().invert_yaxis()
plt.title("Top 15 Feature Importances (Tuned RF)")
plt.show()
```



10 10. Final comparison table

[96]: # 10. Final comparison

```
def evaluate(name, y_true, y_pred):
    print(f"\n{name}")
    print("Accuracy : ", accuracy_score(y_true, y_pred))
    print("Precision:", precision_score(y_true, y_pred, zero_division=0))
    print("Recall   : ", recall_score(y_true, y_pred, zero_division=0))
    print("F1 Score : ", f1_score(y_true, y_pred, zero_division=0))

evaluate("Logistic Regression", y_test, y_pred_log)
evaluate("Decision Tree", y_test, y_pred_tree)
evaluate("Random Forest (baseline)", y_test, y_pred_rf)
evaluate("Random Forest (tuned)", y_test, y_pred_best)
```

Logistic Regression
 Accuracy : 0.7757274662881476
 Precision: 0.5714285714285714
 Recall : 0.6203208556149733
 F1 Score : 0.5948717948717949

Decision Tree
Accuracy : 0.7735982966643009
Precision: 0.569620253164557
Recall : 0.6016042780748663
F1 Score : 0.5851755526657998

Random Forest (baseline)
Accuracy : 0.7927608232789212
Precision: 0.6084656084656085
Recall : 0.6149732620320856
F1 Score : 0.6117021276595744

Random Forest (tuned)
Accuracy : 0.7927608232789212
Precision: 0.6084656084656085
Recall : 0.6149732620320856
F1 Score : 0.6117021276595744